



ARM ARCHITECTURE AND INSTRUCTION SET

EL-GY 6483 Real Time Embedded Systems

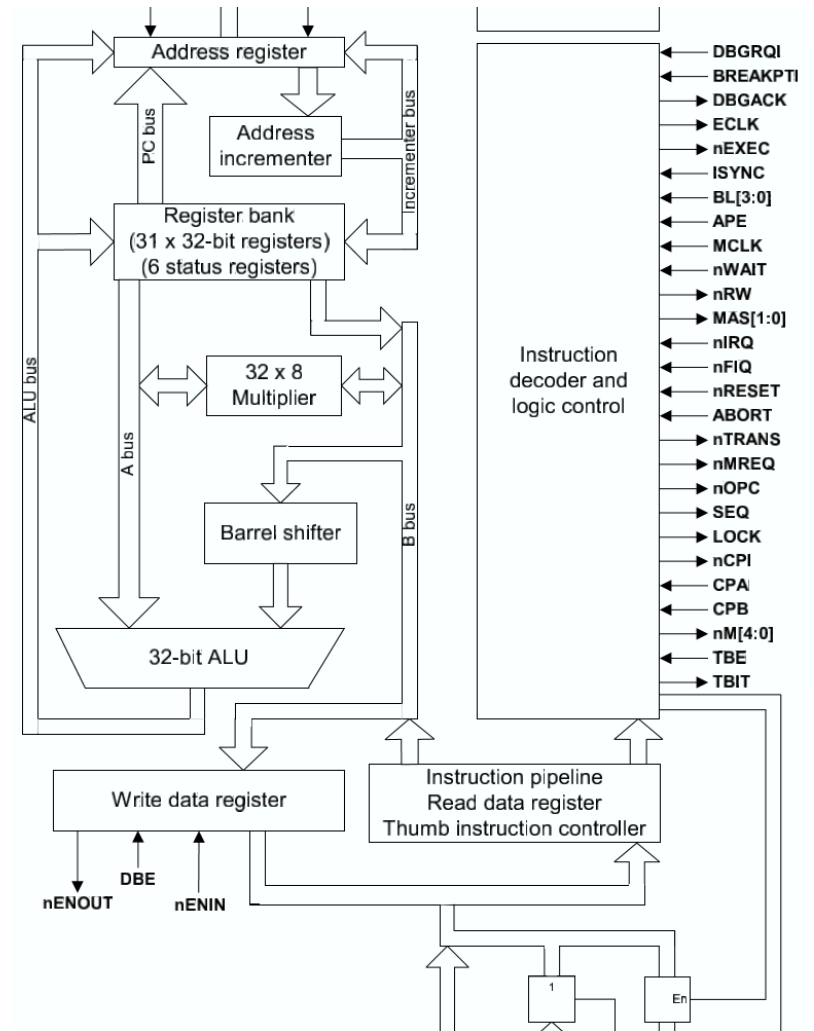
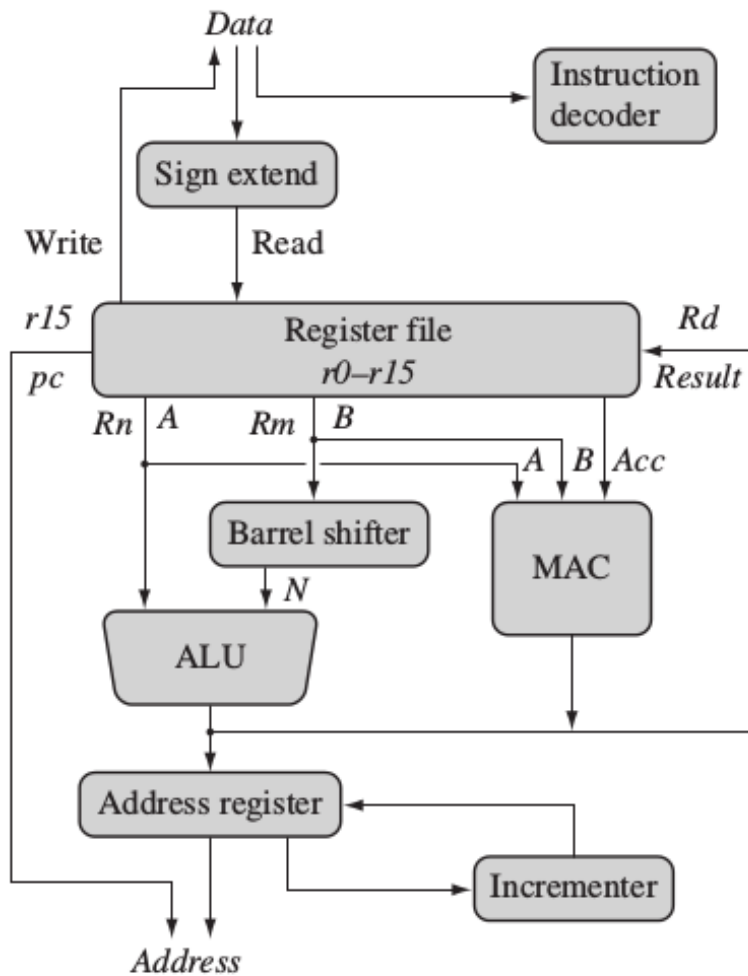


MAIN FEATURES

- 32-bit instructions
- Most executed in single cycle
- Most can be conditionally executed
- Load/store

Also has a compressed 16-bit instruction set (Thumb)

DATAFLOW



PROCESSOR MODES

Six+ operating modes, each with own registers

- Most tasks run in **user** mode
- **FIQ** – entered on high-priority interrupt
- **IRQ** – entered on a normal interrupt
- **Supervisor** (reset), **Abort** (memory access violations), **Undef** (undefined instructions)
- **System** mode added in ARMv4, gives privileged access to same registers as user mode

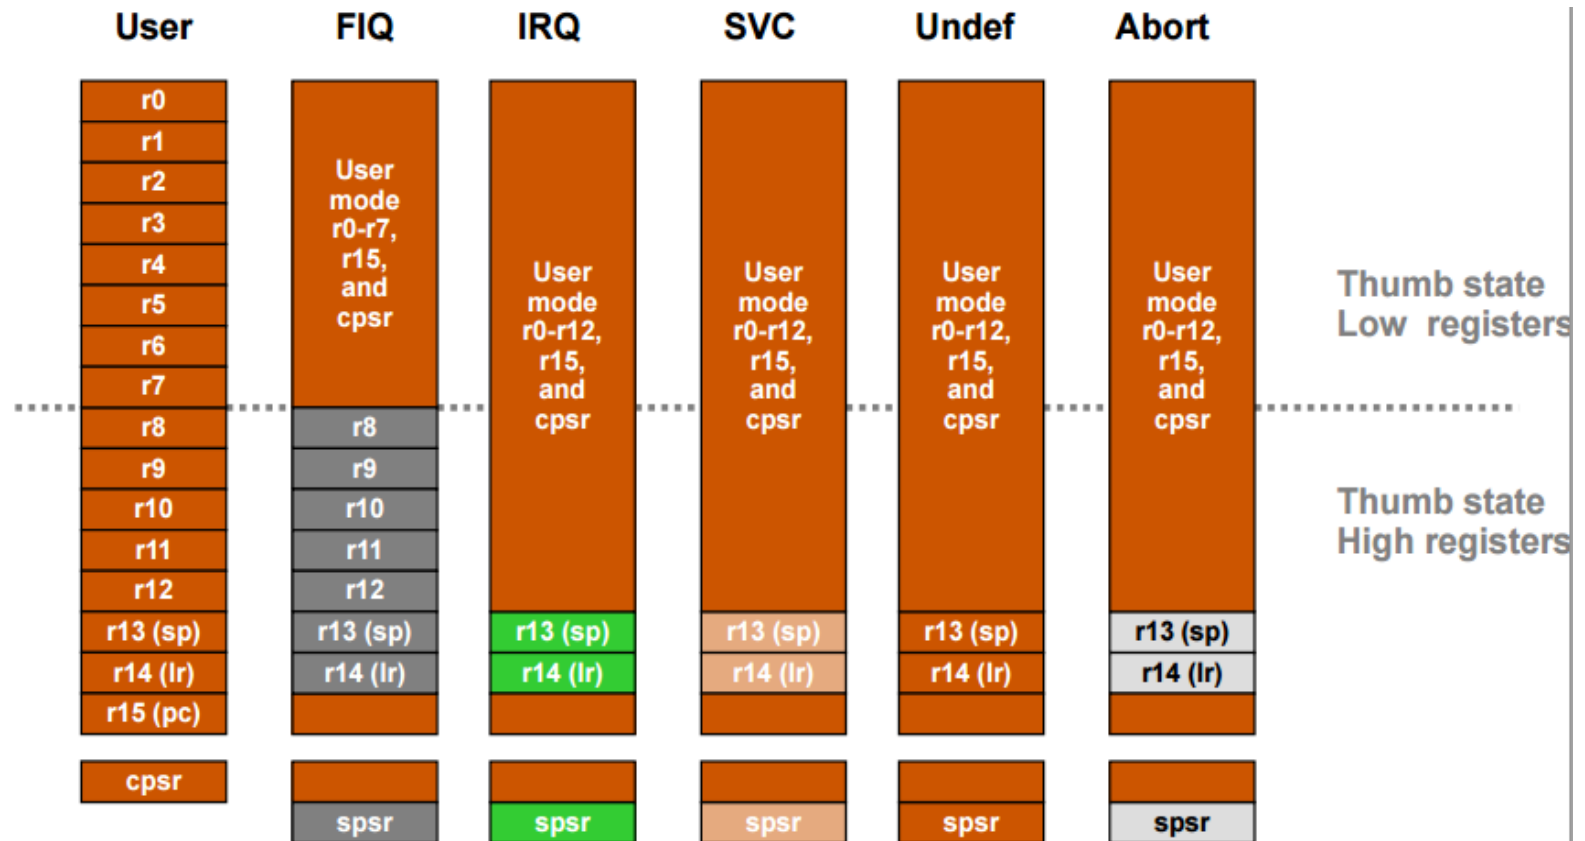
REGISTERS

- 37 registers, each 32 bits long
- Each **mode** can access:
 - A specific set of r0-r12 registers
 - A specific r13 (SP) and r14 (LR)
 - R15 (PC)
 - CPSR
- Privileged modes can also access a SPSR
- Banked registers?

Thumb

- Thumb is a 16-bit instruction set
 - Optimized for code density from C code
 - Improved performance form narrow memory
 - Subset of the functionality of the ARM instruction set
- Core has two execution states –ARM and Thumb
 - Switch between them using BX instruction
- Thumb has characteristic features:
 - Most Thumb instruction are executed unconditionally
 - Many Thumb data process instruction use a 2-address format
 - Thumb instruction formats are less regular than ARM instruction formats, as a result of the dense encoding.

REGISTERS



Note: System mode uses the User mode register set

Processor Modes

- The ARM has six operating modes:
 - User (unprivileged mode under which most tasks run)
 - FIQ (entered when a high priority (fast) interrupt is raised)
 - IRQ (entered when a low priority (normal) interrupt is raised)
 - Supervisor (entered on reset and when a Software Interrupt instruction is executed)
 - Abort (used to handle memory access violations)
 - Undef (used to handle undefined instructions)
- ARM Architecture Version 4 adds a seventh mode:
 - System (privileged mode using the same registers as user mode)

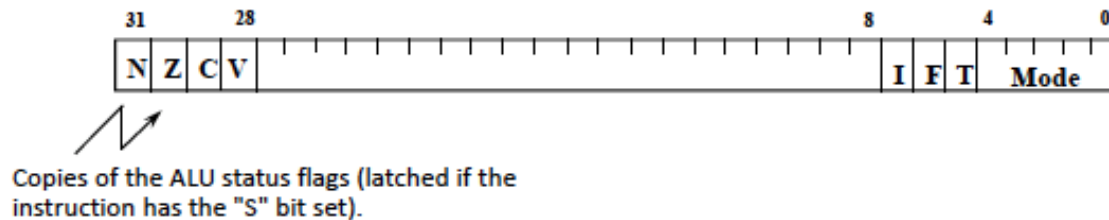
Accessing Registers using ARM Instructions

- All instructions can access r0-r14 directly.
- Most instructions also allow use of the PC.
- Specific instructions to allow access to CPSR and SPSR.
- Note : When in a privileged mode, it is also possible to load-store the (banked out) user mode registers to or from memory

The Program Counter (R15)

- When the processor is executing in ARM state:
 - All instructions are 32 bits in length
 - All instructions must be word aligned
 - Therefore the PC value is stored in bits [31:2] with bits [1:0] equal to zero (as instruction cannot be halfword or byte aligned).
- R14 is used as the subroutine link register (LR) and stores the return address when Branch with Link operations are performed, calculated from the PC.
- Thus to return from a linked branch:
 MOV r15, r14
 or
 MOV pc, lr

The Program Status Registers (CPSR and SPSR)



- Condition Code Flags**

N = **N**egative result from ALU flag

Z = **Z**ero Result from ALU flag

C = ALU operation **C**arried out

V = ALU operation **o**Verflowed

- Interrupt Disable Bits**

I = 1, disables the IRQ

F = 1, disables the FIQ

- Mode Bits**

M[4:0] define the processor mode.

- T bit (Architecture v4T only)**

T = 0, Processor in ARM state

T = 1, Processor in Thumb state

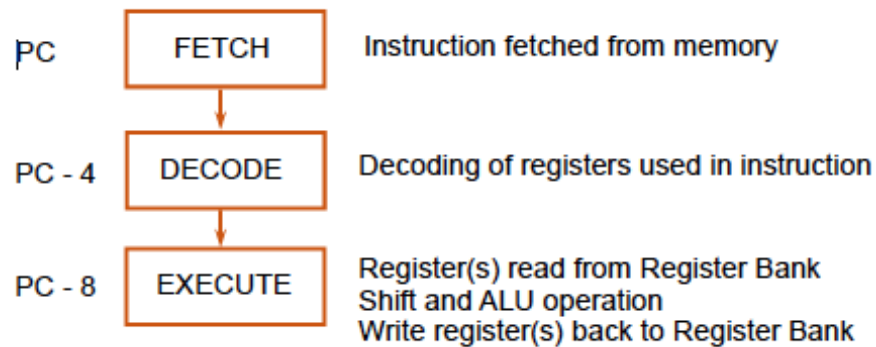
Condition Flags

	Logical Instruction	Arithmetic Instruction
<u>Flag</u>		
Negative (N='1')	No meaning	Bit 31 of the result has been set Indicates a negative number in signed operations
Zero (Z='1')	Result is all zeroes	Result of operation was zero
Carry (C='1')	After Shift operation '1' was left in carry flag	Result was greater than 32 bits
oVerflow (V='1')	No meaning	Result was greater than 31 bits Indicates a possible corruption of the sign bit in signed numbers

Updated by explicit comparison instructions (eg. CMP) and those that use the optional S to specify if the condition code flags must be updated (e.g. ADDS).

ARM Pipeline

- The ARM uses a pipeline in order to increase the speed of the flow of instructions to the processor.
 - Allows several operations to be undertaken simultaneously, rather than serially.



- Rather than pointing to the instruction being executed, the PC points to the instruction being fetched.

INSTRUCTION SYNTAX

`<operation> cond flags Rd,Rn,Operand2`

Basic syntax followed by many instructions. From left to right:

- A three-letter mnemonic, e.g. MOV or ADD.
- An optional two-letter condition code, e.g. EQ or CS.
- An optional additional flag
- The destination register
- First operand register
- Second (more flexible) operand register

ARM Instruction Set Format

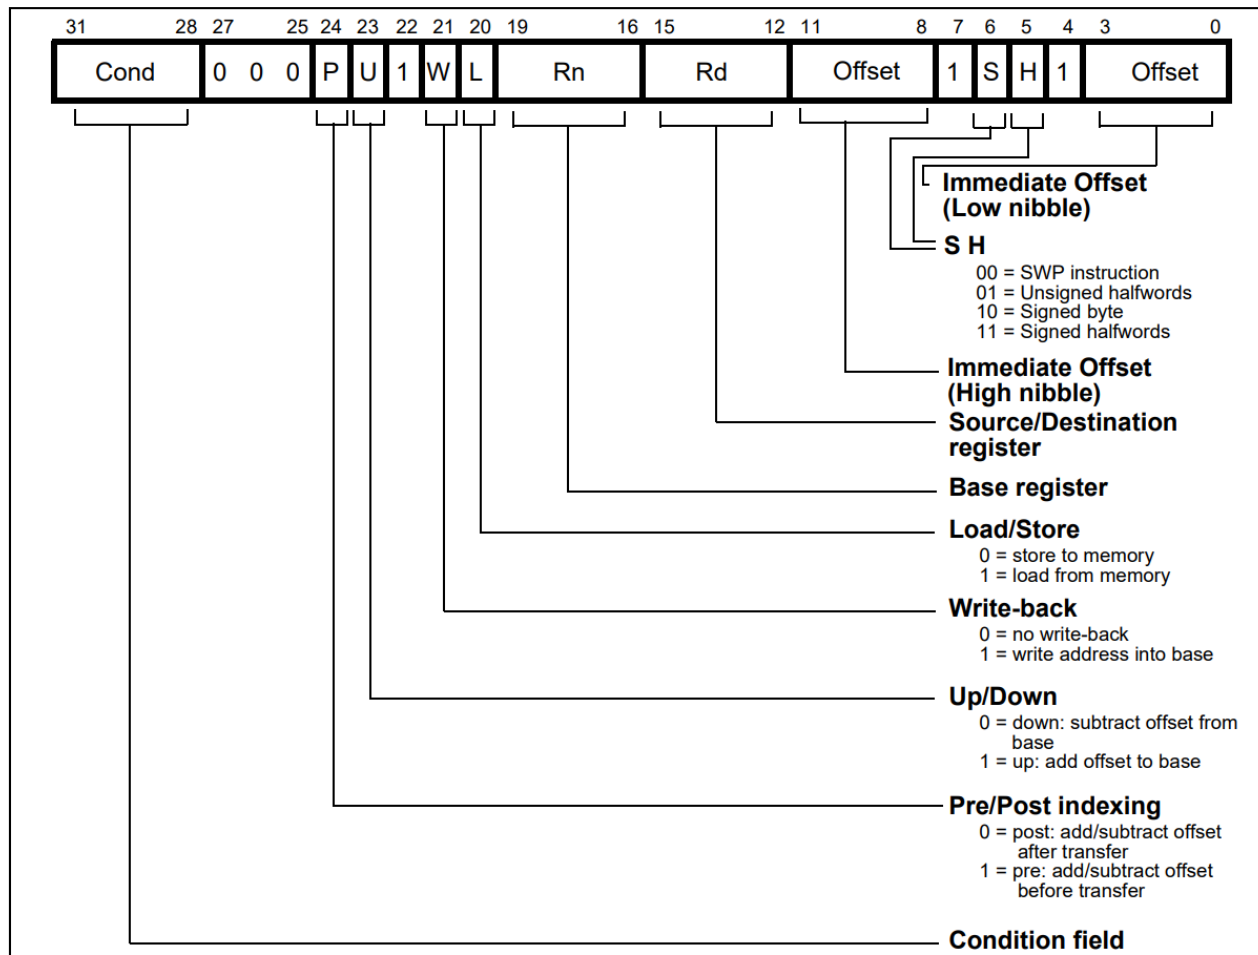
3 1	3 0	2 9	2 8	2 7	2 6	2 5	2 4	2 3	2 2	2 1	2 0	1 9	1 8	1 7	1 6	1 5	1 4	1 3	1 2	1 1	1 0	9	8	7	6	5	4	3	2	1	0	Instruction Type		
Condition				0	0	I	OPCODE				S	Rn				Rs				OPERAND-2												Data processing		
Condition				0	0	0	0	0	0	A	S	Rd				Rn				Rs				1	0	0	1	Rm				Multiply		
Condition				0	0	0	0	1	U	A	S	Rd HIGH				Rd LOW				Rs				1	0	0	1	Rm				Long Multiply		
Condition				0	0	0	1	0	B	0	0	Rn				Rd				0	0	0	0	1	0	0	1	Rm				Swap		
Condition				0	1	I	P	U	B	W	L	Rn				Rd				OFFSET												Load/Store - Byte/Word		
Condition				1	0	0	P	U	B	W	L	Rn				REGISTER LIST																Load/Store Multiple		
Condition				0	0	0	P	U	1	W	L	Rn				Rd				OFFSET 1				1	S	H	1	OFFSET 2				Halfword Transfer Imm Off		
Condition				0	0	0	P	U	0	W	L	Rn				Rd				0	0	0	0	1	S	H	1	Rm				Halfword Transfer Reg Off		
Condition				1	0	1	L	BRANCH OFFSET																								Branch		
Condition				0	0	0	1	0	0	1	0	1	1	1	1	1	1	1	1	1	1	1	1	0	0	0	1	Rn				Branch Exchange		
Condition				1	1	0	P	U	N	W	L	Rn				CRd				CPNum				OFFSET								COPROCESSOR DATA XFER		
Condition				1	1	1	0	Op-1				CRn				CRd				CPNum				OP-2				0	CRm				COPROCESSOR DATA OP	
Condition								OP-1				L	CRn				Rd				CPNum				OP-2				1	CRm				COPROCESSOR REG XFER
Condition				1	1	1	1	SWI NUMBER																								Software Interrupt		

ARM Instruction Set Format

ARM Condition Codes

Opcode [31:28]	Mnemonic extension	Interpretation	Status flag state for execution
0000	EQ	Equal / equals zero	Z set
0001	NE	Not equal	Z clear
0010	CS/HS	Carry set / unsigned higher or same	C set
0011	CC/LO	Carry clear / unsigned lower	C clear
0100	MI	Minus / negative	N set
0101	PL	Plus / positive or zero	N clear
0110	VS	Overflow	V set
0111	VC	No overflow	V clear
1000	HI	Unsigned higher	C set and Z clear
1001	LS	Unsigned lower or same	C clear or Z set
1010	GE	Signed greater than or equal	N equals V
1011	LT	Signed less than	N is not equal to V
1100	GT	Signed greater than	Z clear and N equals V
1101	LE	Signed less than or equal	Z set or N is not equal to V
1110	AL	Always	any
1111	NV	Never (do not use!)	none

ARM Instruction Set Format



COMMON OPERATIONS

- Load and store instructions: LDR, STR, etc.
- Move instructions: MOV, etc.
- Branch (jump) instructions: B, BL, BX, BLX, etc.
- Stack push and pop instructions: PUSH, POP, etc.
- Arithmetic operations: ADD, SUB, MUL, SDIV, UDIV, etc.
- No operation (null operation): NOP

MOVE OPERATIONS

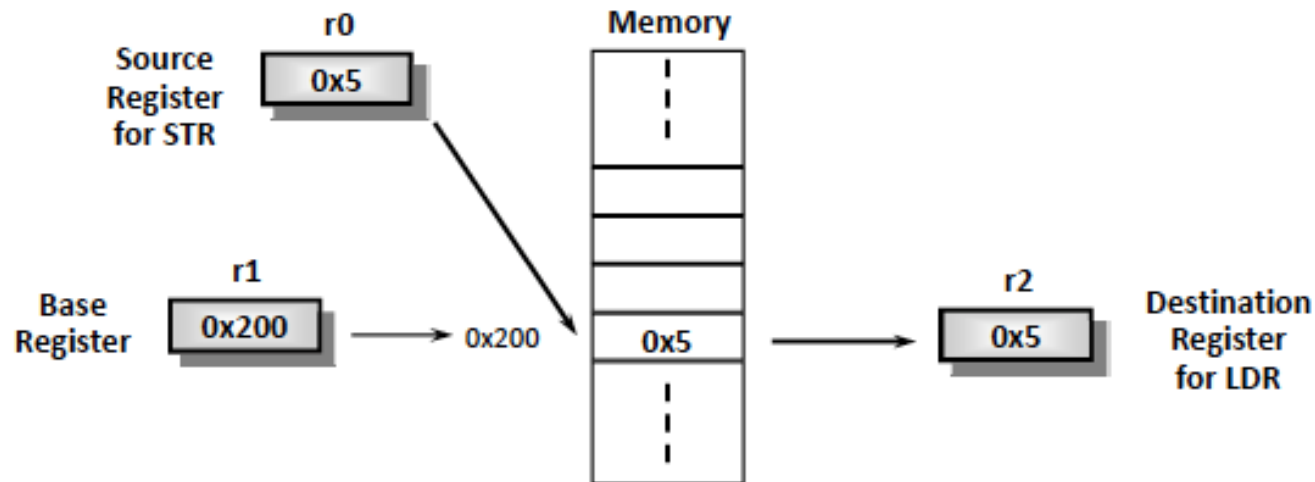
MOV, MVN: “Move” and “Move NOT”; move not does a bitwise logical NOT operation before copying the value to the destination register.

- MOV R0, R1; copy the contents of R1 to R0
- MOV R0, #10; set R0 to 10
- MVN R0, R1; set R0 to bitwise NOT of contents of R1
- MVN R0, #1; set R0 to 0xFFFFFFFF (bitwise NOT of x1)

Load and Store Word or Byte: Base Register

The memory location to be accessed is held in a base register

- `STR r0, [r1]` ; Store contents of r0 to location pointed to
- ; by contents of r1.
- `LDR r2, [r1]` ; Load r2 with contents of memory location
- ; pointed to by contents of r1.

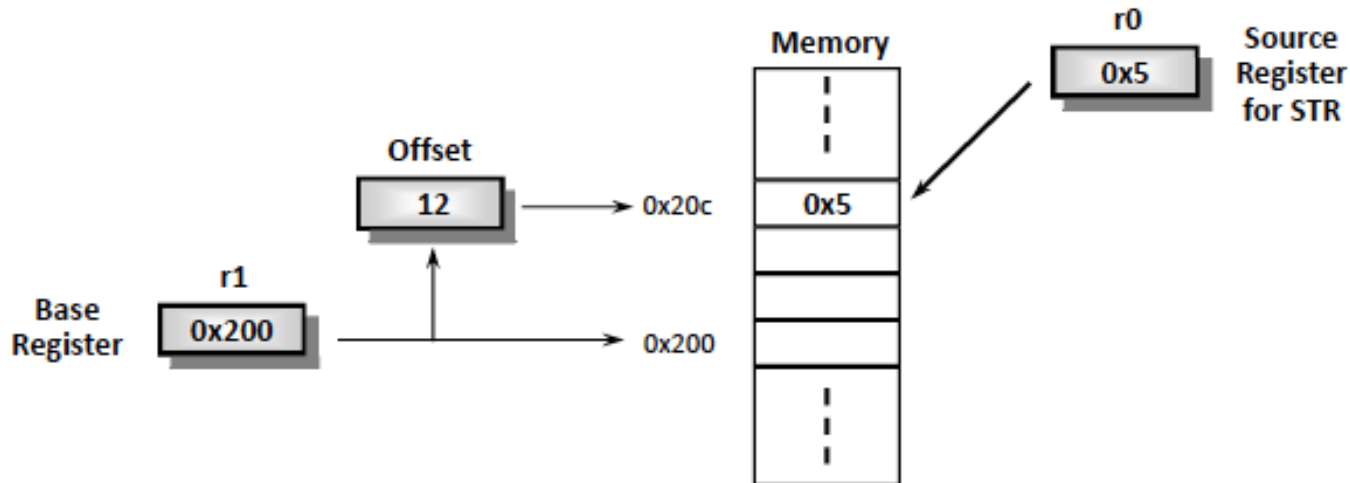


Load/Store: Offsets from the Base Register

- Accessing a location offset from the base register pointer.
- Either:
 - An unsigned 12bit immediate value (ie 0 -4095 bytes).
 - A register, optionally shifted by an immediate value
- This can be either added or subtracted from the base register:
 - Prefix the offset value or register with '+' (default) or '-'.
- This offset can be applied:
 - before the transfer is made: *Pre-indexed addressing*
 - optionally auto-incrementing the base register, by postfixing the instruction with an '!'.
• after the transfer is made: *Post-indexed addressing* causing the base register to be auto-incremented.

Example: Load/Store : Pre-indexed Addressing

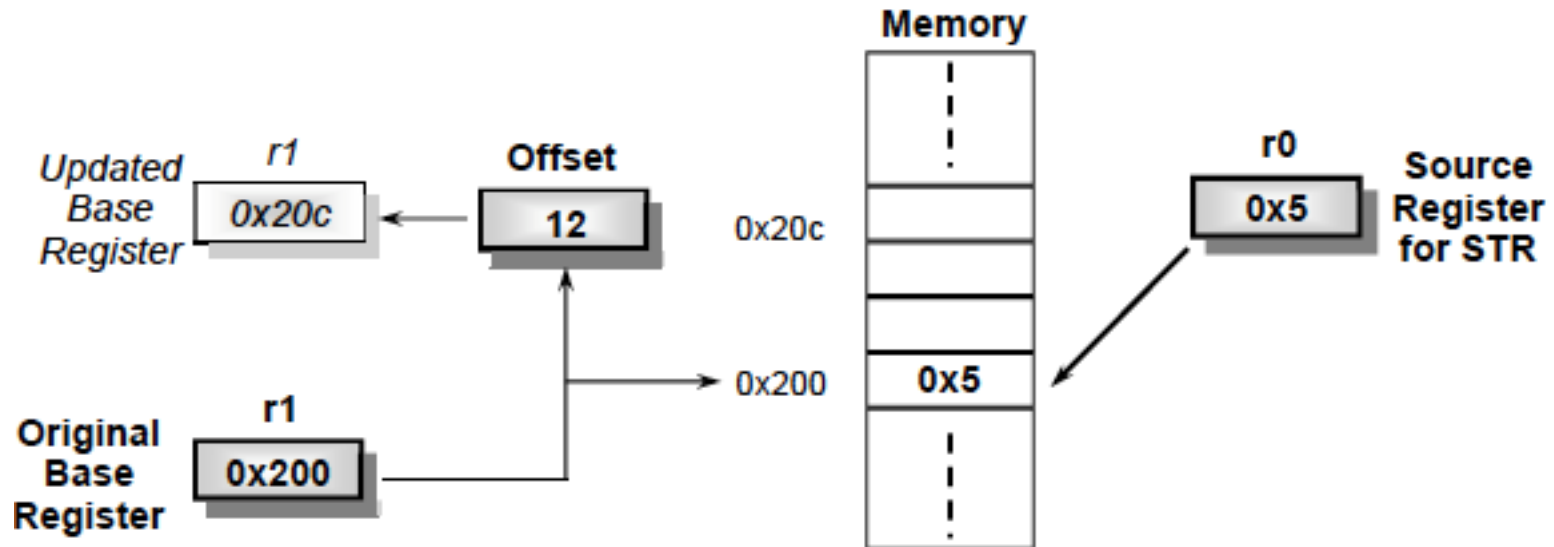
Example: STR r0, [r1,#12]



- To store to location `0x1f4` instead use: `STR r0, [r1,#-12]`
- To auto-increment base pointer to `0x20c` use: `STR r0, [r1, #12]!`
- If `r2` contains 3, access `0x20c` by multiplying this by 4:
`STR r0, [r1, r2, LSL #2]`

Example: Load and Store: Post-indexed Addressing

Example: STR r0, [r1], #12



- To auto-increment the base register to location 0x1f4 instead use:
 - STR r0, [r1], #-12
- If r2 contains 3, auto-increment base register to 0x20c by multiplying this by 4:
 - STR r0, [r1], r2, LSL #2

LOAD AND STORE INSTRUCTIONS

- `LDR R1, [R0]` ; load into R1 the content of the memory location whose address is in R0
- `STR R1, [R0]` ; store the contents of R1 into memory location whose address is in R0
- `LDR R1, [R0, #4]` ; #4 specifies an offset value, load into R1 the content of the memory location whose address is given by the value $R0 + 4$
- `LDR R1, [R0, #4]!` ; #4 specifies an offset value, increment R0 by 4, load into R1 the content of the memory location whose address is given by the new contents of R0
- What about: `STR R1, [R0], #4`

Register names can be in upper case or lower case. Anything after `;` is a comment.

The Barrel Shifter

- The ARM doesn't have actual shift instructions.
- Instead it has a barrel shifter which provides a mechanism to carry out shifts as part of other instructions.
- So what operations does the barrel shifter support?

Barrel Shifter -Left Shift

- Shifts left by the specified amount (multiplies by powers of two)

e.g.

LSL #5 => multiply by 32

Logical Shift Left (LSL)

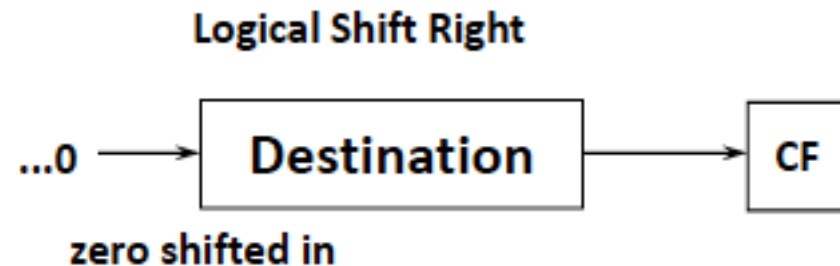


LSR, ASR

Logical Shift Right (LSR)

Shifts right by the specified amount
(divides by powers of two) e.g.

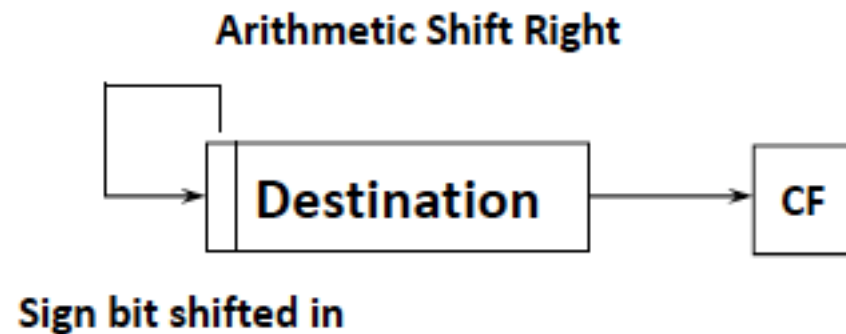
LSR #5 = divide by 32



Arithmetic Shift Right (ASR)

Shifts right (divides by powers of two)
and preserves the sign bit, for 2's
complement operations. e.g.

ASR #5 = divide by 32



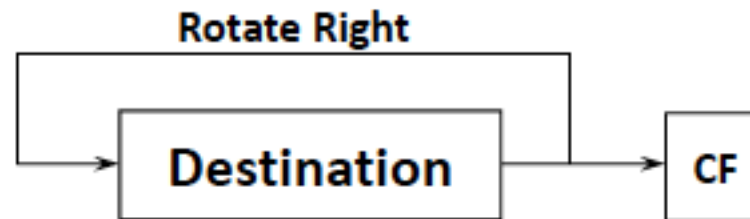
Rotations

Rotate Right (ROR)

Similar to an ASR but the bits wrap around as they leave the LSB and appear as the MSB.

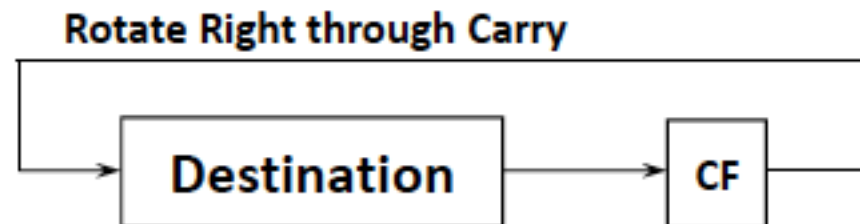
e.g. ROR #5

Note the last bit rotated is also used as the Carry Out.

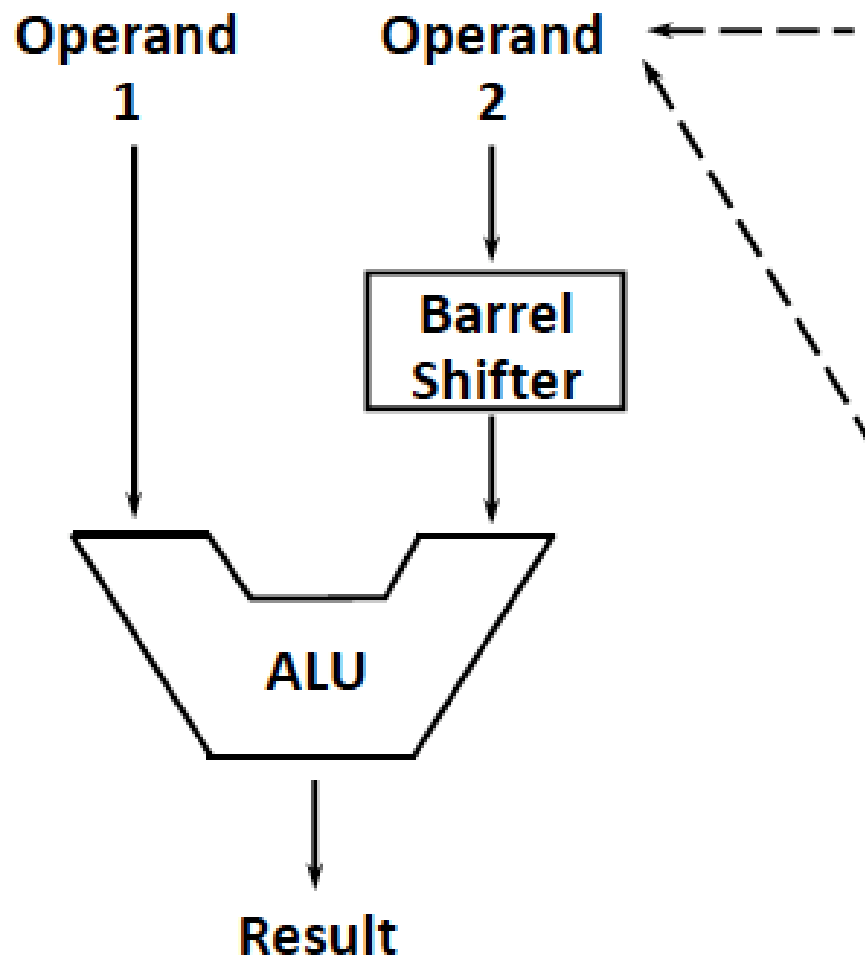


Rotate Right Extended (RRX)

This operation uses the CPSR C flag as a 33rd bit.



Using the Barrel Shifter: The Second Operand



Register, optionally with shift operation applied.

Shift value can be either be:
5 bit unsigned integer
Specified in bottom byte of another register.

****Immediate value**

8 bit number

Can be rotated right through an even number of positions.

Assembler will calculate rotate for you from constant.

BRANCH OPERATIONS

- B and BL are branch with “immediate” arguments (e.g. jumping to a label/function); the symbol . Is a synonym for the current program location
- BX and BLX are branch with “register” arguments (i.e. jump to an address stored in a register)
- BL and BLX store a bookmark to the current place in the program by writing the address of the next instruction in the link register (LR).
 - B labelA; branch to the label labelA
 - BL labelA; update LR and branch to the label labelA
 - BX LR; branch to the location whose address is in LR, e.g. return from a function call
 - BLX R0; update LR and branch to the location whose address is in R0
 - B . ; branch to the current program location (infinite loop)

PUSH AND POP OPERATIONS (STACK)

`PUSH {R1}` ; push the contents of R1 onto the stack

`PUSH {R0,R1}` ; push the contents of R0 and R1 onto the stack

`PUSH {R0,R2-R4}` ; push the contents of R0, R2, R3, R4 onto the stack

`PUSH {R0,LR}` ; push the contents of R0, LR onto the stack

`POP {R0,R1}` ; pop the top two 32-bit values from the stack into R0 and R1

PUSH and POP operations update the SP.

ADDITION AND SUBTRACTION

ADD , SUB : Add and subtract

ADC , SBC : “Add with carry” and “Subtract with carry”

The carry instructions also utilize the carry flag (set by previous instructions – the carry flag is stored as a bit in the application program status register)

- ADD R1, R0, R1 ; set R1 to the sum of the contents of R0 and R1
- SUB R2, R0, R1 ; set R2 to the difference of the contents of R0 and R1
- ADD R2, R0, #10 ; $R2 = R0 + 10$

RSB : reverse subtract (i.e., subtract the contents of the second operand from the third operand), e.g.,

- RSB R2, R0, R1 ; set R2 to the difference of the contents of R1 and R0

Signed and unsigned variants: SADD16, etc.

MULTIPLICATION AND DIVISION

MUL, MLA, and MLS : “Multiply”, “Multiply with accumulate”, and “Multiply with subtract”

- MUL R2, R0, R1 ; $R2 = R0 * R1$
- MLA R3, R0, R1, R2 ; $R3 = R0 * R1 + R2$
- MLS R3, R0, R1, R2 ; $R3 = R0 * R1 - R2$

Signed and unsigned variants: SMLA, etc. SDIV, UDIV : “Signed division” and “Unsigned division”

- SDIV R2, R0, R1 ; signed divide, $R2 = R0 / R1$
- UDIV R2, R0, R1 ; unsigned divide, $R2 = R0 / R1$

FLOATING POINT INSTRUCTIONS

Floating point instructions are available if there is a floating point unit (FPU) in the system and is enabled (the FPU is generally enabled as part of the start-up sequence on ARM Cortex-M4F).

- VADD, VSUB, etc. : floating point addition and subtraction
- VMUL, VDIV, etc. : floating point multiplication and division
- VABS : floating point absolute value

BITWISE OPERATIONS

AND, ORR (logical OR), etc.; exclusive OR is EOR

LSL, LSR : “logical shift left” and “logical shift right”

ROR : rotate right

- AND R2, R0, R1 ; $R2 = R0 \& R1$
- OR R2, R0, R1 ; $R2 = R0 | R1$
- AND R2, R0, #0x10 ; $R2 = R0 \& 0x10$
- OR R2, R0, #0x10 ; $R2 = R0 | 0x10$
- LSL R2, R0, #2 ; $R2 = R0 \ll 2$

CONDITIONAL EXECUTION OF INSTRUCTIONS

Many instructions support optional suffixes denoting various conditions to specify that the instruction must be executed only if the specified condition is true.

- Some examples of condition suffixes:
 - EQ for equal (i.e., $Z = 1$), NE for not equal (i.e., $Z = 0$)
 - GT for “greater than” and LT for “less than” ; these conditions are evaluated using combinations of flags Z, N, and V. For example, GT is equivalent to “ $Z = 0$ and $N = V$ ”, LT is equivalent to “ $N \neq V$ ”
 - GE and LE for “greater than or equal” and “less than or equal”
- Example: ADDEQ will do an addition only if the “equal” condition is currently active (i.e., a previous instruction caused the Z flag to become 1); BNE will do a branch (jump) only if the “equal” condition is not currently active.

- In effect, a SWI is a user-defined instruction.
- It causes an exception trap to the SWI hardware vector (thus causing a change to supervisor mode, plus the associated state saving), thus causing the SWI exception handler to be called.
- The handler can then examine the comment field of the instruction to decide what operation has been requested.
- By making use of the SWI mechanism, an operating system can implement a set of privileged operations which applications running in user mode can request.
- See Exception Handling Module for further details.

- In effect, a SWI is a user-defined instruction.
- It causes an exception trap to the SWI hardware vector (thus causing a change to supervisor mode, plus the associated state saving), thus causing the SWI exception handler to be called.
- The handler can then examine the comment field of the instruction to decide what operation has been requested.
- By making use of the SWI mechanism, an operating system can implement a set of privileged operations which applications running in user mode can request.
- See Exception Handling Module for further details.

Assembler: Pseudo-ops

- AREA -> chunks of data (\$data) or code (\$code)
- ADR -> load address into a register
- ADR R0, BUFFER
- ALIGN -> adjust location counter to word boundary usually after a storage directive
- END -> no more to assemble

Assembler: Pseudo-ops

- Example:

```
AREA  cacheable, CODE, ALIGN=3
```

```
rout1  ; code      ; aligned on 8-byte boundary
```

```
        ; code
```

```
        MOV    pc,lr ; aligned only on 4-byte boundary
```

```
        ALIGN  8     ; now aligned on 8-byte boundary
```

```
rout2  ; code
```

Assembler: Pseudo-ops

- `IMPORT` -> name of routine to import for use in this routine
 - `IMPORT _printf ; C print routine`
- `EXPORT` -> name of routine to export for use in other routines
 - `EXPORT add2 ; add2 routine`
- `EQU` -> symbol replacement
 - `loopcnt EQU 5`

Assembly Line Format

- *label*<whitespace> *instruction*<whitespace> ; *comment*
- *label*: created by programmer, alphanumeric
- whitespace: space(s) or tab character(s)
- *instruction*: op-code mnemonic or pseudo-op with required fields
- *comment*: preceded by ; ignored by assembler but useful to the programmer for documentation
- NOTE: All fields are optional.

REFERENCES

- Sloss, Andrew, Dominic Symes, and Chris Wright. "ARM system developer's guide: designing and optimizing system software." Morgan Kaufmann, 2004.
- ARM University Program. "ARM Processors and Architectures Comprehensive Overview." <http://arm.com>, 2012.